

A Symplectic N-Body Solver for Solar System Dynamics: Validation Against JPL Horizons

Karl Keshavarzi^{*1}

¹Department of Electrical & Computer Engineering
University of Waterloo

February 2026

Abstract

We present a high-performance N-body gravitational simulator for solar system dynamics, integrating 35 bodies (the Sun, 8 planets, Pluto, and 25 natural satellites) over timescales of up to 249 years. The simulator employs a 4th-order Yoshida symplectic integrator, achieving energy conservation of $\Delta E/E \leq 7.4 \times 10^{-11}\%$ over 100 years and $\leq 1.0 \times 10^{-10}\%$ over 249 years with a timestep of $\Delta t = 360\text{ s}$. Positional accuracy is validated against the NASA JPL Horizons DE441 ephemeris, yielding a mean maximum relative error of 0.05% over 100 years across all 35 bodies. Residual errors are attributed to physics model differences; principally, the absence of post-Newtonian relativistic corrections, solar oblateness J_2 , and asteroid perturbations, rather than numerical integration error. The implementation features OpenMP parallelization with SIMD vectorization for large- N scalability, a Structure-of-Arrays memory layout for cache performance, and an automated validation pipeline interfacing directly with the JPL Horizons API.

Keywords: N-body simulation, symplectic integration, Yoshida integrator, solar system dynamics, JPL Horizons, gravitational dynamics

^{*}This work was completed independently and is not affiliated with the University of Waterloo.

1. Introduction

The gravitational N-body problem, that is, computing the trajectories of N mutually interacting masses, is one of the oldest and most fundamental problems in computational physics. Accurate long-duration integration of planetary orbits requires numerical methods that preserve the geometric structure of Hamilton's equations. Non-symplectic integrators exhibit secular energy drift that corrupts orbital elements over extended timescales [2].

This work presents a complete N-body simulation framework for solar system dynamics. The system integrates 35 bodies (the Sun, 8 planets, Pluto, and 25 natural satellites) under pairwise Newtonian gravitation using the 4th-order Yoshida symplectic integrator [1]. We validate the simulator against the NASA JPL Horizons system, which provides state vectors computed from the DE441 planetary ephemeris. DE441 is a comprehensive model incorporating post-Newtonian relativity, solar oblateness, asteroid perturbations from over 300 bodies, and tidal effects [3].

The primary contributions of this work are:

1. Implementation and verification of a 4th-order Yoshida symplectic integrator achieving $\Delta E/E \leq 10^{-10}\%$ over multi-century timescales.
2. Comprehensive positional validation against JPL DE441 for 35 solar system bodies, with quanti-

tative attribution of residual errors to specific missing physics.

3. An automated validation pipeline that fetches JPL Horizons ephemerides and performs epoch-aligned comparison, eliminating interpolation artifacts.

2. Theoretical Background

2.1. Equations of Motion

The gravitational acceleration of body i due to all other $N - 1$ bodies is given by Newton's law of universal gravitation:

$$\mathbf{a}_i = \sum_{\substack{j=1 \\ j \neq i}}^N \frac{G m_j (\mathbf{r}_j - \mathbf{r}_i)}{(|\mathbf{r}_j - \mathbf{r}_i|^2 + \varepsilon^2)^{3/2}} \quad (1)$$

where $G = 6.6743 \times 10^{-11} \text{ m}^3 \text{ kg}^{-1} \text{ s}^{-2}$ is the gravitational constant, m_j is the mass of body j , and \mathbf{r}_i and \mathbf{r}_j are position vectors. A softening parameter $\varepsilon = 10^{-9} \text{ m}$ is included to prevent numerical singularity at zero separation. This value is negligible relative to interplanetary distances and does not measurably affect the dynamics; it serves only as a numerical safeguard. The results are insensitive to the choice of ε provided $\varepsilon \ll \min_{i \neq j} |\mathbf{r}_i - \mathbf{r}_j|$.

2.2. Symplectic Integration

Hamiltonian systems possess a symplectic structure: the phase-space flow preserves the canonical 2-form $\sum_i dp_i \wedge dq_i$. Standard Runge-Kutta methods violate this structure, causing secular energy drift proportional to integration time. Symplectic integrators, by construction, exactly preserve a nearby shadow Hamiltonian $\tilde{H} = H + \mathcal{O}(\Delta t^p)$, bounding energy error for all time [2].

2.2.1. Velocity Verlet 2nd-Order Integrator

The Velocity Verlet algorithm [4] is a 2nd-order symplectic method widely used in molecular dynamics:

$$\mathbf{r}(t + \Delta t) = \mathbf{r}(t) + \Delta t \mathbf{v}(t) + \frac{1}{2} \Delta t^2 \mathbf{a}(t) \quad (2)$$

$$\mathbf{v}(t + \Delta t) = \mathbf{v}(t) + \frac{1}{2} \Delta t [\mathbf{a}(t) + \mathbf{a}(t + \Delta t)] \quad (3)$$

This requires storing the previous acceleration and one force evaluation per timestep.

2.2.2. Yoshida 4th-Order Integrator

The Yoshida method [1] constructs a 4th-order symplectic scheme by composing three leapfrog sub-steps with carefully chosen coefficients. Define:

$$w_1 = \frac{1}{2 - 2^{1/3}}, \quad w_0 = -\frac{2^{1/3}}{2 - 2^{1/3}} \quad (4)$$

The position drift and velocity kick coefficients are:

$$c_1 = c_4 = \frac{w_1}{2}, \quad d_1 = d_3 = w_1 \quad (5)$$

$$c_2 = c_3 = \frac{w_0 + w_1}{2}, \quad d_2 = w_0 \quad (6)$$

Each timestep proceeds as four position drifts interleaved with three force evaluations ("kicks"):

$$\mathbf{r} \xrightarrow{c_1} \mathbf{v} \xrightarrow{d_1} \mathbf{r} \xrightarrow{c_2} \mathbf{v} \xrightarrow{d_2} \mathbf{r} \xrightarrow{c_3} \mathbf{v} \xrightarrow{d_3} \mathbf{r} \xrightarrow{c_4} \quad (7)$$

Here $\mathbf{r} \xrightarrow{c_k}$ denotes $\mathbf{r} \leftarrow \mathbf{r} + c_k \Delta t \mathbf{v}$ and $\mathbf{v} \xrightarrow{d_k}$ denotes $\mathbf{v} \leftarrow \mathbf{v} + d_k \Delta t \mathbf{a}(\mathbf{r})$.

The scheme achieves $\mathcal{O}(\Delta t^4)$ local truncation error with only 3 force evaluations per step, compared to 4 for classical RK4. The negative coefficient $w_0 < 0$ (equivalently $d_2 < 0$) introduces a backward sub-step that cancels lower-order error terms. The symplectic structure guarantees that energy errors remain bounded and oscillatory, never exhibiting secular growth.

3. Implementation

3.1. Software Architecture

The simulator is implemented in C++17 with a modular object-oriented architecture. An abstract **Integrator** base class defines the integration interface, with **Yoshida** and **Velocity_Verlet** as concrete implementations. An abstract **Force** base class permits extensible force models; the current implementation provides **Gravity** for pairwise Newtonian interactions. The **Simulation** class orchestrates the

time-stepping loop, energy diagnostics, and CSV output.

All simulation parameters (timestep Δt , duration, and output interval) are specified in a single `Config.hpp` header as `constexpr` values. This enables compile-time computation of derived quantities such as total steps and output cadence. A `static_assert` enforces that the output interval is an exact multiple of the timestep.

3.2. Data Layout: Structure of Arrays

Particle data is stored in a Structure-of-Arrays (SoA) layout rather than the conventional Array-of-Structures (AoS). Each physical attribute (position, velocity, acceleration, and mass) is stored in a contiguous heap-allocated array:

```
pos_x[0..N-1], pos_y[0..N-1], ...
```

This layout maximizes cache-line utilization during the inner force loop. The computation accesses all x -positions contiguously, then all y -positions, and so on. On modern architectures with 64-byte cache lines, this yields 8 contiguous `double` values per cache-line fetch. In contrast, AoS layout scatters a single particle's fields across multiple cache lines, degrading spatial locality.

Raw array pointers are qualified with compiler-specific `_restrict_` annotations, informing the compiler that no aliasing occurs between position, velocity, and acceleration arrays. This enables aggressive SIMD auto-vectorization of the inner loops.

3.3. Force Computation

The gravitational force calculation uses $\mathcal{O}(N^2)$ direct pairwise summation. Self-interaction ($i = j$) is eliminated with a branchless floating-point mask:

$$\text{mask}_j = \begin{cases} 0.0 & \text{if } i = j \\ 1.0 & \text{otherwise} \end{cases} \quad (8)$$

The force contribution is multiplied by this mask, avoiding a conditional branch that would disrupt

SIMD vectorization. The inner loop is annotated with `#pragma omp simd` with reductions on the acceleration accumulators to enable vectorized execution. For $N \geq 500$, the outer loop is parallelized with OpenMP via `#pragma omp parallel for`.

Note that Newton's third law symmetry ($\mathbf{F}_{ij} = -\mathbf{F}_{ji}$) is not exploited; each body's acceleration is computed independently by summing over all N others. This doubles the floating-point work relative to a symmetric $j > i$ loop but preserves the regular access pattern required for SIMD vectorization and avoids race conditions under thread parallelism.

3.4. Initial Conditions

State vectors for all 35 bodies are sourced from the NASA JPL Horizons system at epoch J1950.0 (January 1, 1950, 00:00 UTC). The reference frame is solar system barycentric (SSB) with the ecliptic reference plane (ICRF). Positions are given in kilometers and velocities in km s^{-1} , converted to SI units at initialization. The body catalog comprises the Sun, 8 planets (Mercury through Neptune), Pluto, and 25 natural satellites: the Moon, Phobos, Deimos, the four Galilean moons, Amalthea, 9 Saturnian moons (Mimas through Phoebe), 5 Uranian moons (Ariel through Miranda), Triton, Nereid, and Charon.

The simulation operates in the SSB frame throughout. Because all initial state vectors are sourced from JPL Horizons in this frame, the total momentum of the 35-body subset is not identically zero (the missing solar system bodies carry the complementary momentum). No barycentric correction is applied.

An automated Python pipeline (`jpl_COMPARE.py fetch`) queries the Horizons API, retrieves state vectors at the configured output cadence, and generates both the C++ initial conditions header and the reference ephemeris for validation.

3.5. Build System

The project uses CMake (minimum version 3.16) for cross-platform builds. Compiler optimizations (`-O3 -march=native`) and OpenMP linkage are configured

automatically. The build targets GCC and Clang via MinGW Makefiles on Windows or native Makefiles on Unix systems.

4. Validation Methodology

Validation is performed by direct comparison of simulated trajectories against JPL Horizons ephemeris data. The pipeline ensures methodological rigor through several design choices.

Epoch-aligned comparison. Both the simulator and the JPL fetch use identical output intervals (`output_hours`), sourced from the same `Config.hpp` file. This eliminates interpolation error in the comparison.

Relative position error. For each body at each epoch, the relative position error is defined as:

$$\epsilon_{\text{rel}}(t) = \frac{|\mathbf{r}_{\text{sim}}(t) - \mathbf{r}_{\text{JPL}}(t)|}{|\mathbf{r}_{\text{JPL}}(t)|} \quad (9)$$

where $|\cdot|$ denotes the Euclidean norm. For each body, we report the maximum of $\epsilon_{\text{rel}}(t)$ over all epochs $t > 0$.

Energy conservation. Total mechanical energy $E = T + U$ is sampled at regular intervals throughout the integration. Energy conservation is quantified by the peak-to-peak relative variation:

$$\Delta E/E = \frac{|E_{\max} - E_{\min}|}{|E_0|} \quad (10)$$

where E_{\max} and E_{\min} are the extrema over all sampled epochs, and E_0 is the initial energy. Note that this peak-to-peak measure can differ by up to a factor of 2 from the more commonly reported metric $\max_t |E(t) - E_0|/|E_0|$ [2, 7]. The peak-to-peak definition provides a conservative upper bound on the energy variation.

5. Results

5.1. 100-Year Simulation

Table 1 summarizes results for a 100-year integration (1950–2050) with $\Delta t = 360$ s and the Yoshida

integrator. The simulation completes in 56.5 s on a single core.

Table 1: 100-year simulation summary (35 bodies, $\Delta t = 360$ s, Yoshida 4th-order).

Metric	Value
Mean max. relative error	0.053%
Worst-case body	Ganymede (0.290%)
Planet error range	0.0003%–0.173%
Moon error range	0.009%–0.290%
Energy conservation $\Delta E/E$	$7.4 \times 10^{-11}\%$
Wall-clock runtime	56.5 s

Table 2 presents the per-body breakdown, grouped by dynamical class.

Table 2: Per-body maximum relative position error (%) for the 100-year simulation.

Body	Error (%)	Body	Error (%)
<i>Sun and Planets</i>			
Saturn	0.0003	Venus	0.022
Uranus	0.0003	Sun	0.071
Neptune	0.0004	Mercury	0.173
Jupiter	0.0009		
Mars	0.006		
Pluto	0.010		
Earth	0.013		
<i>Selected Moons</i>			
Iapetus	0.009	Io	0.114
Miranda	0.010	Titan	0.125
Phobos	0.014	Europa	0.175
Triton	0.016	Callisto	0.188
Moon	0.018	Ganymede	0.290

5.2. 249-Year Simulation

Extending the integration to 249 years (1950–2199) tests the long-term stability of the symplectic integrator. Table 3 compares the two simulation durations, and Table 4 presents the per-body breakdown.

Table 3: Comparison of 100-year and 249-year simulations ($\Delta t = 360$ s).

Metric	100 yr	249 yr
Mean max. error	0.053%	0.115%
Worst body	Ganymede (0.29%)	Sun (1.34%)
Worst planet	Mercury (0.17%)	Mercury (0.42%)
$\Delta E/E$	$7.4 \times 10^{-11}\%$	$1.0 \times 10^{-10}\%$
Runtime	56.5 s	171.6 s

Table 4: Per-body maximum relative position error (%) for the 249-year simulation.

Body	Error (%)	Body	Error (%)
<i>Sun and Planets</i>			
Saturn	0.0007	Earth	0.031
Uranus	0.0007	Venus	0.055
Neptune	0.0009	Mercury	0.417
Jupiter	0.002	Sun	1.336
Mars	0.015		
Pluto	0.034		
<i>Selected Moons</i>			
Miranda	0.010	Io	0.115
Triton	0.016	Nereid	0.149
Umbriel	0.019	Europa	0.183
Iapetus	0.022	Titan	0.184
Phobos	0.022	Hyperion	0.226
Moon	0.045	Ganymede	0.290
Amalthea	0.050	Callisto	0.411

Several observations follow from these results.

Energy conservation remains bounded. The energy drift increases only from $7.4 \times 10^{-11}\%$ to $1.0 \times 10^{-10}\%$ when the integration duration grows by a factor of 2.5. This sub-linear growth is characteristic of symplectic integrators, which preserve a shadow Hamiltonian \tilde{H} and exhibit only oscillatory energy variations, never secular drift.

Positional errors grow sub-linearly. The mean error roughly doubles (0.053% to 0.115%) over $2.5\times$ the integration time. This is consistent with slow phase drift accumulation, not exponential divergence. The outer planets (Saturn, Uranus, Neptune) remain below $10^{-3}\%$ even at 249 years, reflecting the dominance of solar gravitation in their dynamics and the absence of significant missing physics at their orbital distances.

The Sun becomes the worst-case body at 249 years. This is a coordinate artifact, not a simulation failure. In the solar system barycentric frame, the Sun’s position is determined by the gravitational influence of the planets (predominantly Jupiter and Saturn). The Sun’s barycentric displacement is only $\sim 10^6$ km, roughly one solar radius, while outer planets orbit at $\sim 10^9$ km. A given absolute position error therefore produces a relative error roughly 10^3 – $10^4\times$ larger for the Sun than for a giant planet. At 100 years the Sun’s relative error is 0.07%; by 249 years, accumulated absolute error yields 1.34% relative error. In absolute terms, this remains physically negligible.

Callisto overtakes Ganymede at 249 years. Callisto’s error grows from 0.188% to 0.411%, surpassing Ganymede (0.290% at both durations). This is consistent with the longer orbital period of Callisto (16.7 d vs. 7.15 d), which couples to slightly different phase drift accumulation rates over multi-century timescales. Hyperion (0.226%) also emerges as a notable outlier among Saturnian moons, reflecting its chaotic rotation and irregular orbit.

5.3. Sources of Residual Error

The residual errors are dominated by physics model differences between the simulator (Newtonian gravity, 35 bodies) and JPL DE441. The principal missing contributions are described below.

Post-Newtonian relativity. The Einstein–Infeld–Hoffmann (EIH) equations contribute ~ 43 arcsec century $^{-1}$ of precession to Mercury’s orbit [5]. This is the dominant error source for Mercury (0.17% at 100 yr, 0.42% at 249 yr), as the Newtonian model cannot capture this effect. Implementation of 1PN corrections was attempted; however, the velocity-dependent post-Newtonian potential breaks the separability assumption ($H = T(\mathbf{p}) + V(\mathbf{q})$) required by the Yoshida integrator. This degraded energy conservation by ~ 4 orders of magnitude (from $10^{-11}\%$ to $10^{-7}\%$), producing worse overall accuracy than the purely Newtonian model. Implicit

symplectic splitting schemes [6] or mixed-variable symplectic methods [5] offer potential paths to incorporating these corrections without sacrificing symplecticity.

Solar oblateness (J_2). The Sun’s quadrupole moment produces a small additional precession, particularly affecting Mercury’s perihelion advance.

Asteroid perturbations. JPL DE441 includes gravitational contributions from over 300 asteroids (principally Ceres, Pallas, and Vesta). These bodies are absent from the 35-body model.

Moon phase drift. Satellites with short orbital periods accumulate phase errors over many orbits. Io orbits Jupiter every 1.77 d and completes $\sim 20,600$ orbits in 100 years. At $\Delta t = 360$ s, each orbit is resolved by ~ 425 timesteps. This is adequate for 4th-order per-orbit accuracy, but small per-orbit phase errors compound over thousands of orbits. Reducing Δt or implementing adaptive timestepping would improve satellite accuracy at the cost of runtime.

Tidal and non-gravitational effects. Tidal dissipation, solar radiation pressure, and planetary ring interactions are included in DE441 but absent from the simulator.

5.4. Confirmation: Integration Error Is Not the Limiting Factor

The claim that residual errors are physics-limited rather than numerics-limited is supported by two observations.

First, reducing Δt from 3600 s to 360 s (a 10 \times reduction) improves energy conservation by a factor of $\sim 460\times$ (from $4.8 \times 10^{-8}\%$ to $1.0 \times 10^{-10}\%$). For a 4th-order method, a 10 \times reduction in Δt is expected to improve the energy error by 10 $^4\times$. The observed 460 \times improvement indicates that the energy error has not yet fully converged at $\Delta t = 3600$ s but is approaching the asymptotic regime. Meanwhile, positional errors against JPL remain essentially unchanged, having converged to the physics model floor.

Second, the error hierarchy matches the expected

importance of missing physics: Mercury (relativity) > Galilean moons (phase drift) > outer planets (minimal missing effects).

6. Performance

6.1. Computational Complexity

The force computation scales as $\mathcal{O}(N^2)$ per timestep due to direct pairwise summation. The Yoshida integrator requires 3 force evaluations per step, yielding a total cost of $\mathcal{O}(3N^2 \cdot T/\Delta t)$ for a simulation of duration T .

For the 35-body solar system with $\Delta t = 360$ s over 100 years, this corresponds to $3 \times 35^2 \times 8,766,900 \approx 3.2 \times 10^{10}$ pairwise interaction evaluations, completed in 56.5 s on a single core.

6.2. Parallelization

The outer loop of the force computation is parallelized with OpenMP for particle counts exceeding a configurable threshold ($N \geq 500$). The inner loop uses `#pragma omp simd` with reductions to enable SIMD auto-vectorization. For the 35-body solar system, the overhead of thread management exceeds the parallelization benefit, so OpenMP threading is automatically disabled. SIMD vectorization of the inner loop remains active regardless of N .

To validate the parallelization at scale, the solar system body catalog was replicated to produce a 10,000-body test case. On a 6-core CPU, OpenMP threading yielded a $\sim 6\times$ speedup over the single-threaded baseline, demonstrating near-linear scaling. This is expected: the $\mathcal{O}(N^2)$ pairwise force loop is embarrassingly parallel over the outer index, and at $N = 10,000$ the computation-to-synchronization ratio is large enough to saturate all cores.

6.3. Memory Efficiency

The SoA layout stores 10 scalar fields per particle (3 positions, 3 velocities, 3 accelerations, and 1 mass) plus 3 additional acceleration fields retained for the Velocity Verlet integrator, totaling $13N$ doubles or $104N$ bytes. For $N = 35$, the entire particle

dataset occupies ~ 3.6 KB, fitting comfortably within L1 cache (typically 32–64 KB). The `__restrict__` qualifiers permit the compiler to assume no pointer aliasing, enabling store-to-load forwarding optimizations.

7. Conclusion

We have presented a symplectic N-body simulator that achieves near-machine-precision energy conservation ($\Delta E/E \leq 10^{-10}\%$) over multi-century integration of the solar system. Positional accuracy against JPL Horizons DE441 is 0.05% mean maximum relative error over 100 years across 35 bodies. The Yoshida 4th-order integrator provides an excellent balance between accuracy and computational cost, requiring only 3 force evaluations per timestep while delivering 4th-order convergence and guaranteed bounded energy error.

The residual positional errors are conclusively attributed to physics model limitations rather than numerical integration error. This is demonstrated by the convergence of positional accuracy with decreasing timestep, while energy conservation continues to improve by orders of magnitude.

Future work includes implementing post-Newtonian corrections via an implicit or mixed-variable symplectic scheme, adding solar oblateness perturbations, incorporating adaptive timestepping for improved resolution of fast-orbiting satellites, and performing a systematic convergence study over a range of timesteps.

Code Availability

The source code for the simulator and the validation pipeline are open source and available on GitHub.

References

- [1] H. Yoshida, “Construction of higher order symplectic integrators,” *Phys. Lett. A*, vol. 150, no. 5–7, pp. 262–268, 1990.
- [2] E. Hairer, C. Lubich, and G. Wanner, *Geometric Numerical Integration: Structure-Preserving Algorithms for Ordinary Differential Equations*, 2nd ed. Berlin: Springer, 2006.
- [3] R. S. Park *et al.*, “The JPL planetary and lunar ephemerides DE440 and DE441,” *Astron. J.*, vol. 161, no. 3, p. 105, 2021.
- [4] L. Verlet, “Computer ‘experiments’ on classical fluids. I. Thermodynamical properties of Lennard-Jones molecules,” *Phys. Rev.*, vol. 159, no. 1, pp. 98–103, 1967.
- [5] J. Wisdom and M. Holman, “Symplectic maps for the N -body problem,” *Astron. J.*, vol. 102, pp. 1528–1538, 1991.
- [6] J. E. Chambers, “A hybrid symplectic integrator that permits close encounters between massive bodies,” *Mon. Not. R. Astron. Soc.*, vol. 304, no. 4, pp. 793–799, 1999.
- [7] H. Rein and D. S. Spiegel, “IAS15: a fast, adaptive, high-order integrator for gravitational dynamics, accurate to machine precision over a billion orbits,” *Mon. Not. R. Astron. Soc.*, vol. 446, no. 2, pp. 1424–1437, 2015.